



TICKET REPETITION ANALYZER


Finding the Same Problem, Reported Differently

How AI-based duplicate detection cut through 11,700+ support tickets

Project Review · STAR Summary · Prepared for CTO & Engineering Leadership

Same Issue, Different Words — and No Way to See It

- Support teams handle 11,000+ tickets across products and branches
- The same underlying problem is reported repeatedly in different wording
- Keyword search and manual tagging both miss these matches
- Result: no visibility into what's actually recurring, or how much effort is duplicated



Same problem — invisible to keyword search


The cost: repeat effort is invisible — teams keep re-solving problems that were already fixed, with no data to prove it or prioritize it.

TASK

What We Set Out to Build

- Automatically detect tickets describing the same issue, regardless of wording
- Quantify how many tickets are duplicates, and how much of it is still unresolved
- Deliver it as a self-serve dashboard — readable by anyone, no engineering help needed
- Keep all ticket data on-premise — no sensitive data sent to an external service
- Make it work from day one on a real export — no long integration project required

How It Works, Technically




Similarity threshold is tunable — lower it to catch more duplicates, raise it to avoid false matches.

Why it works: two tickets don't need matching keywords — only a similar meaning — so "VPN not connecting" and "unable to access VPN" land in the same group automatically.

What We Delivered

- Upload-and-go: drag in a MongoDB export, no server setup
- 5 interactive dashboard views, switchable in one click
- Auto-recommended view — opens on whatever fits this data best, in plain English
- Runs fully on-premise: nothing leaves the machine
- ~11,700 tickets processed in about 1 minute

The dashboard, 5 views:



The Impact, in Numbers

11,692

Tickets analyzed

947

Recurring issues found

34%

Repeat rate

2,985

Redundant tickets

191

Still open right now

- 1 in 3 tickets handled today is a repeat of a problem already reported
- 191 open tickets could potentially close the moment their original issue is fixed
- Team leads can now prioritize by actual repeat volume — not guesswork